


# Predictive Modeling Prototyping with the R caret Package

Prof L



# About Me

Clinical Assistant Professor at Purdue Universities' Krannert School of Management and Co-Founder/Chief Data Scientist of Biz Analytics Lab, LLC in Lafayette, IN.

- Teach and mentor students (Fall/Spring semesters)
- Work with a couple fantastic partners (Summer)

At Krannert, course coordinator and teacher for:

- MGMT 571 Data Mining (Fall semester)
- MGMT 590 Using R for Analytics (Fall semester)
- MGMT 590 Predictive Analytics (Spring semester)
- MGMT 690 MS BAIM Industry Practicum (Spring semester)

Spend most of my time obtaining and mentoring experimental learning projects for students within Purdue's M.S. in Business Analytics & Information Management (BAIM) program.

# Key items to take away

- How **caret** is designed to work
- Key **caret** functions
- Purdue's MS BAIM is awesome!

# Presentation and R script link

<https://github.com/MatthewALanham/informsba2018>



# Data Mining (MS BAIM fall core)

- Following a process (e.g. CRISP-DM, INFORMS CAP framework)
- Relational Databases/SQL
- Supervised vs. Unsupervised Learning
- Regression vs. Classification Problems
- Cross-validation Designs (validation set, k-fold, LOOCV, bias-variance tradeoff)
- Exploratory Data Analysis & Visualization with Tableau
- Data Pre-Processing (multicollinearity, binning, feature engineering)
- Linear Models
- Dimension Reduction via PCA & Stepwise Approaches
- Clustering Analysis (Hierarchical, k-Means, PAM, SOMs, Silhouettes)
- Classification and Regression Trees
- Feed Forward Neural Networks
- Association Rules/Market Basket Analysis


\* Can easily do using caret

# Predictive Analytics (MS BAIM spring elective)

- INFORMS CAP framework
- Decision Model Basics (LP, IP, MIP)
- Designing Solutions & Integrating Analytics (Descriptive, Predictive, Prescriptive)
- Review [Cross-validation Designs](#), Bias-variance Tradeoff
- [KNN](#), [Bayes Classifier](#), [Naïve Bayes](#)
- [Linear & Quadratic Discriminant Analysis \(LDA/QDA\)](#)
- [Support Vector Machines](#), Factorization Machines
- [Multi-classification Modeling & Evaluation](#) - Multinomial Logit, SVM
- Cost-based Learning & Evaluation
- Ensembling (Voting, Propensity Averaging, [Random Forest](#)/Bagging, [AdaBoost](#)/Boosting, [Gradient Boosting Machines](#), Meta-Modeling)
- [Recurrent, LSTM, & Convolution Neural Nets](#) to Deep learning
- Text Mining

\* Can easily do using caret

# R/RStudio continues to be popular in practice




Source: KDnuggets.com (2017)



Source: G2crowd.com (2018)

# MS BAIM students use other software too...



Source: KDnuggets.com (2017)




Source: G2crowd.com (2018)



# caret Package

The **C**lassification **A**nd **RE**gression **T**raining **caret** package was developed by Max Kuhn as a tool to streamline the predictive modeling process for R users.

His book published in 2013 and has many nice examples.



Additional information here:

<http://topepo.github.io/caret/index.html>

Most of the functionality can be broken down into 5 areas

Data splitting

Pre-processing

Feature selection

Model tuning  
using resampling

Variable  
importance  
estimation

# Example: Classic Adult/Census Income dataset

**Problem type:** Binary Classification

**Objective:** Predict whether income exceeds \$50K/yr for an individual based on 1994 census database.

**Dataset:** <http://archive.ics.uci.edu/ml/datasets/Adult>

Data Set Characteristics:	Multivariate	Number of Instances:	48842	Area:	Social
Attribute Characteristics:	Categorical, Integer	Number of Attributes:	14	Date Donated	1996-05-01
Associated Tasks:	Classification	Missing Values?	Yes	Number of Web Hits:	1142995

# Features

```
# Features:
#age: continuous.
#workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov
#           , State-gov, Without-pay, Never-worked.
#fnlwgt: continuous.
#education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm
#           , Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate
#           , 5th-6th, Preschool.
#education-num: continuous.
#marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed
#               , Married-spouse-absent, Married-AF-spouse.
#occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial
#            , Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical
#            , Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv
#            , Armed-Forces.
#relationship: wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
#race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
#sex: Female, Male.
#capital-gain: continuous.
#capital-loss: continuous.
#hours-per-week: continuous.
#native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany
#               , Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China
#               , Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietn
#               , Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ec
#               , Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland
#               , Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong
#               , Holand-Netherlands.
#income: >50K, <=50K.
```

Potential features

Target

# Load dataset

```
# Load data from the web
myUrl <- "http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data"
d <- read.table(file=myUrl, header=F, sep=" ", quote="",
               colClasses=c("numeric", "factor", "numeric", "factor", "numeric",
                           , rep("factor", 5), rep("numeric", 3), rep("factor", 2)))

# specify column names
names(d) <- c("age", "workclass", "fnlwt", "education", "educationnum",
             "maritalstatus", "occupation", "relationship", "race", "sex",
             "capitalgain", "capitalloss", "hoursperweek", "nativecountry",
             "income")

# examine data structure
str(d)
```

RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function Addins

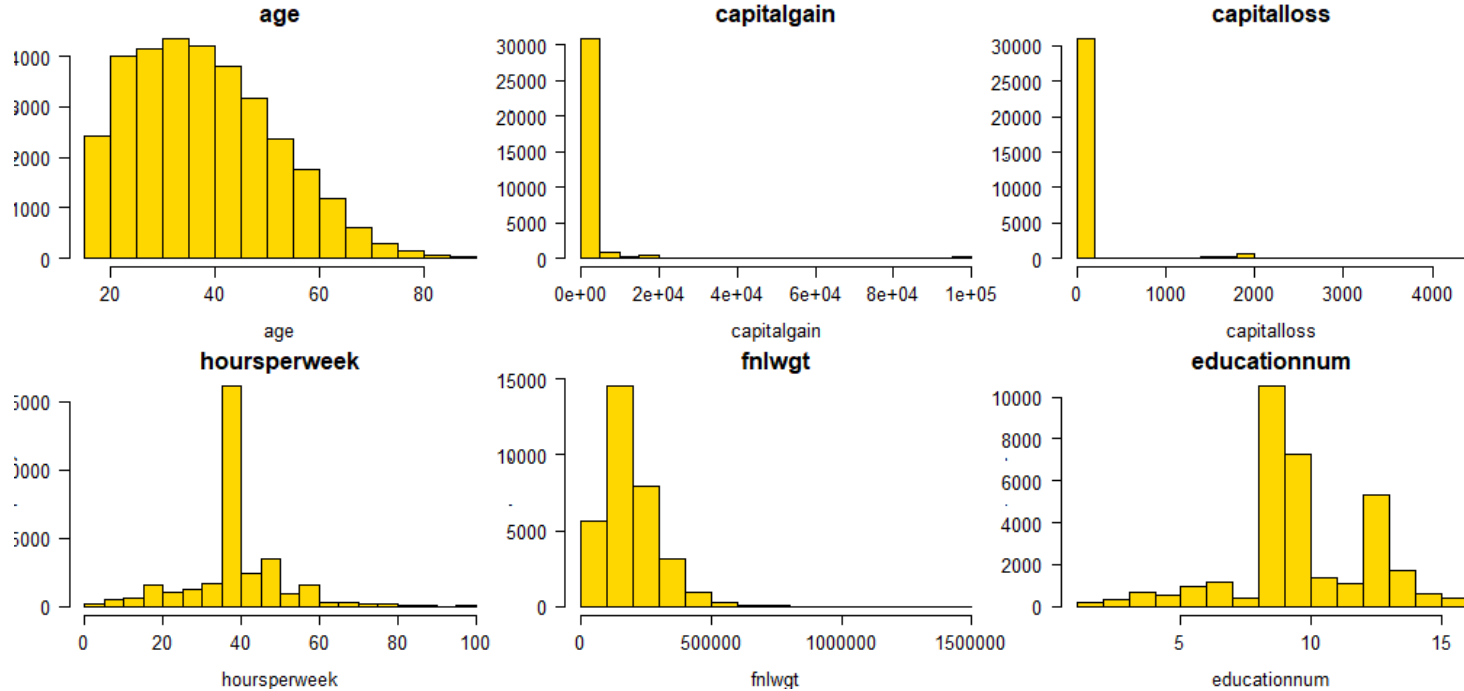
Project: (None)

INFORMScaRet.R.R\* d



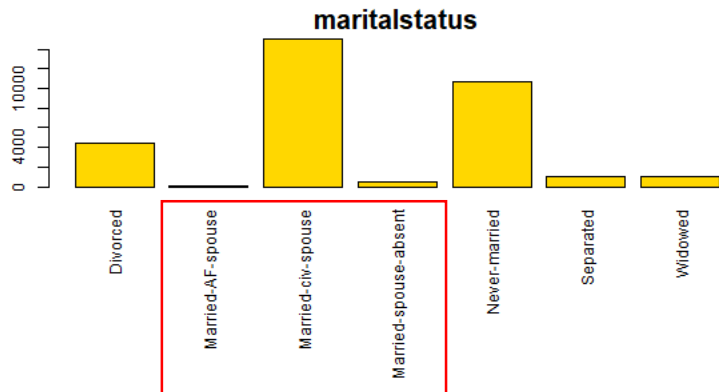
Filter

	age	workclass	fnlwt	education	educationnum	maritalstatus	occupation	relationship	race	sex	capitalgain	capitalloss	hoursperweek	nativecountry	income
1	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
2	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
3	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
4	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
5	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K
6	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States	<=50K





# EDA



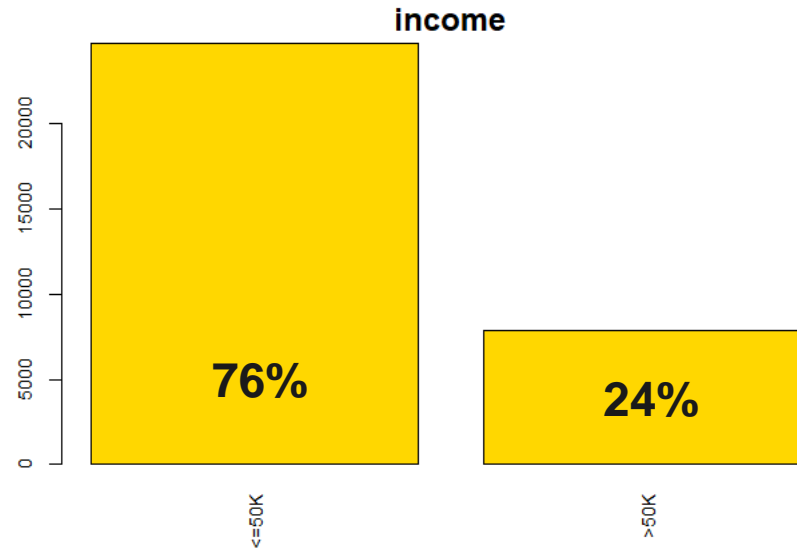
# EDA



# EDA



# EDA





# EDA

- In the R script I do some data cleaning. Will not cover for this presentation.

	y	age	workclass	fnlwgt	education	educationnum	maritalstatus	occupation	relationship	race	sex	capitalgain	capitalloss	hoursperweek	nativecountry
1	<=50K	39	Gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States
2	<=50K	50	Self	83311	Bachelors	13	Married	Exec-managerial	Husband	White	Male	0	0	13	United-States
3	<=50K	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States
4	<=50K	53	Private	234721	11th	7	Married	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States
5	<=50K	28	Private	338409	Bachelors	13	Married	Prof-specialty	Wife	Black	Female	0	0	40	Cuba
6	<=50K	37	Private	284582	Masters	14	Married	Exec-managerial	Wife	White	Female	0	0	40	United-States
7	<=50K	49	Private	160187	9th	5	Married	Other-service	Not-in-family	Black	Female	0	0	16	Jamaica

# caret preprocessing functions

- 1) **dummyVars ()** function creates a full set of dummy variables (i.e. less than full rank parameterization)
- 2) **findCorrelation()** function searches through a correlation matrix and returns a vector of integers corresponding to columns to remove to reduce pair-wise correlations
- 3) **findLinearCombos()** function enumerates and resolves the linear combinations in a numeric matrix
- 4) **preProcess()** function performs transformations (centering, scaling etc.) estimated from the training data and applied to any data set with the same variables

# dummyVars()

- 1) **dummyVars ()** function creates a full set of dummy variables (i.e. less than full rank parameterization)

Data  
d 30162 obs. of 15 variables

y	age	workclass	fnlwgt	education	educati
<=50K	39	Gov	77516	Bachelors	
<=50K	50	Self	83311	Bachelors	
<=50K	38	Private	215646	HS-grad	
<=50K	53	Private	234721	11th	
<=50K	28	Private	338409	Bachelors	



d 30162 obs. of 100 variables

y	age	workclassGov	workclassPrivate	workclassSelf	workclassWithoutpay	fn
<=50K	39	1	0	0	0	
<=50K	50	0	0	1	0	
<=50K	38	0	1	0	0	
<=50K	53	0	1	0	0	
<=50K	28	0	1	0	0	

# dummyVars()


- 1) **dummyVars ()** function creates a full set of dummy variables (i.e. less than full rank parameterization)

```
#####  
## Creating Dummy Variables  
#####  
# Here we want to create a dummy 0/1 variable for every level of a categorical  
# variable  
library(caret)  
dummies <- dummyVars(y ~ ., data = d) # create dummies for Xs  
ex <- data.frame(predict(dummies, newdata = d)) # actually creates the dummies  
names(ex) <- gsub("\\\\.", "", names(ex)) # removes dots from col names  
d <- cbind(d$y, ex) # combine target var with Xs  
names(d)[1] <- "y" # name target var 'y'  
rm(dummies, ex) # clean environment  
#####
```

# findCorrelation()

2) **findCorrelation()** function searches through a correlation matrix and returns a vector of integers corresponding to columns to remove to reduce pair-wise correlations.

If you build a model that has highly correlated independent variables it can lead to unstable models because it will tend to weight those more even though they might not be that important.



# findCorrelation()

```
#####  
# Identify Correlated Predictors and remove them  
#####  
# If you build a model that has highly correlated independent variables it can  
# lead to unstable models because it will tend to weight those more even though  
# they might not be that important  
  
# calculate correlation matrix using Pearson's correlation formula  
descrCor <- cor(d[,2:ncol(d)]) # correlation matrix  
highCorr <- sum(abs(descrCor[upper.tri(descrCor)]) > .85) # num Xs with cor > t  
summary(descrCor[upper.tri(descrCor)]) # summarize the cors  
  
# which columns in your correlation matrix have a correlation greater than some  
# specified absolute cutoff?  
highlyCorDescr <- findCorrelation(descrCor, cutoff = 0.85)  
filteredDescr <- d[,2:ncol(d)][,-highlyCorDescr] # remove those specific columns  
descrCor2 <- cor(filteredDescr) # calculate a new cor matrix  
# summarize those correlations to see if all features are now within our range  
summary(descrCor2[upper.tri(descrCor2)])  
  
# update dataset by removing those filtered variables that were highly correlated  
d <- cbind(d$y, filteredDescr)  
names(d)[1] <- "y"  
  
rm(filteredDescr, descrCor, descrCor2, highCorr, highlyCorDescr) # clean up  
#####
```

# findCorrelation()

I define highly correlated as 0.85 in this example

```
# calculate correlation matrix using Pearson's correlation formula
descrCor <- cor(d[,2:ncol(d)]) # correlation matrix
highCorr <- sum(abs(descrCor[upper.tri(descrCor)]) > .85) # num Xs with cor > t
summary(descrCor[upper.tri(descrCor)]) # summarize the cors
```

highCorr tells me how many variables have a correlation +/- 0.85

```
> highCorr
[1] 2
```

```
> summary(descrCor[upper.tri(descrCor)]) # summarize the cors
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-1.000000	-0.008280	-0.001848	-0.004943	0.003857	0.874673

Looks like we have high correlations

# findCorrelation()

findCorrelation() identifies those columns in the data that are highly correlated

```
# which columns in your correlation matrix have a correlation greater than some  
# specified absolute cutoff?  
highlyCorDescr <- findCorrelation(descrCor, cutoff = 0.85)  
filteredDescr <- d[,2:ncol(d)][,-highlyCorDescr] # remove those specific columns  
descrCor2 <- cor(filteredDescr) # calculate a new cor matrix  
# summarize those correlations to see if all features are now within our range  
summary(descrCor2[upper.tri(descrCor2)])
```

```
> highlyCorDescr  
[1] 43 54
```

```
> summary(descrCor2[upper.tri(descrCor2)])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-0.794808	-0.008033	-0.001881	-0.004856	0.003477	0.505892

All correlations are now below our threshold



# findCorrelation()

Assuming you can justify dropping those features in respect to the problem, here we combine our target variable with the set of features are not “highly correlated.”

```
# update dataset by removing those filtered vars that were highly correlated  
d <- cbind(d$y, filteredDescr)  
names(d)[1] <- "y"  
  
rm(filteredDescr, descrCor, descrCor2, highCorr, highlyCorDescr) # clean up
```

# findLinearCombos()

3) findLinearCombos() function enumerates and resolves the linear combinations in a numeric matrix

Y	Age	workclassGov	workclassPrivate	workclassSelf	workclassWithoutpay	row sum
1	18	1	0	0	0	1
1	22	1	0	0	0	1
0	45	1	0	0	0	1
1	33	0	0	1	0	1
0	56	0	0	1	0	1
1	43	0	1	0	0	1
0	51	0	1	0	0	1
0	25	0	0	0	1	1

Drop a column

Y	Age	workclassGov	workclassPrivate	workclassSelf	workclassWithoutpay	row sum
1	18	1	0	0	0	0
1	22	1	0	0	0	0
0	45	1	0	0	0	0
1	33	0	0	1	0	1
0	56	0	0	1	0	1
1	43	0	1	0	0	1
0	51	0	1	0	0	1
0	25	0	0	0	1	1

# findLinearCombos()

```
#####  
# Identifying linear dependencies and remove them  
#####  
# Find if any linear combinations exist and which column combos they are.  
# Below I add a vector of 1s at the beginning of the dataset. This helps ensure  
# the same features are identified and removed.  
library(caret)  
# first save response  
y <- d$y  
  
# create a column of 1s. This will help identify all the right linear combos  
d <- cbind(rep(1, nrow(d)), d[2:ncol(d)])  
names(d)[1] <- "ones"  
  
# identify the columns that are linear combos  
comboInfo <- findLinearCombos(d)  
comboInfo  
  
# remove columns identified that led to linear combos  
d <- d[, -comboInfo$remove]  
  
# remove the "ones" column in the first column  
d <- d[, c(2:ncol(d))]  
  
# Add the target variable back to our data.frame  
d <- cbind(y, d)  
  
rm(y, comboInfo) # clean up
```

# findLinearCombos()

Here we can set the different groups of columns that form linear combinations

```
> comboInfo <- findLinearCombos(d)
```

```
> comboInfo
```

```
$linearCombos
```

```
$linearCombos[[1]]
```

```
[1] 6 1 3 4 5
```

Work class dummies

```
$linearCombos[[2]]
```

```
[1] 23 1 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22
```

```
$linearCombos[[3]]
```

```
[1] 24 1 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22
```

```
$linearCombos[[4]]
```

```
[1] 29 1 26 27 28
```

Marital status dummies

```
$linearCombos[[5]]
```

```
[1] 43 1 30 31 32 33 34 35 36 37 38 39 40 41 42
```

```
$linearCombos[[6]]
```

```
[1] 49 1 44 45 46 47 48
```

```
$linearCombos[[7]]
```

```
[1] 54 1 50 51 52 53
```

```
$linearCombos[[8]]
```

```
[1] 56 1 55
```

```
$linearCombos[[9]]
```

```
[1] 100 1 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77
```

```
[21] 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97
```

```
[41] 98 99
```

Native country dummies

```
$remove
```

```
[1] 6 23 24 29 43 49 54 56 100
```

# findLinearCombos()

Here we can see the columns to drop. By default it will drop the first column among a set of linear combos.

```
> comboInfo <- findLinearCombos(d)
> comboInfo
$linearCombos
$linearCombos[[1]]
[1] 6 1 3 4 5

$linearCombos[[2]]
[1] 23 1 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22

$linearCombos[[3]]
[1] 24 1 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22

$linearCombos[[4]]
[1] 29 1 25 26 27 28

$linearCombos[[5]]
[1] 43 1 30 31 32 33 34 35 36 37 38 39 40 41 42

$linearCombos[[6]]
[1] 49 1 44 45 46 47 48

$linearCombos[[7]]
[1] 54 1 50 51 52 53

$linearCombos[[8]]
[1] 56 1 55

$linearCombos[[9]]
[1] 100 1 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77
[21] 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97
[41] 98 99

$remove
[1] 6 23 24 29 43 49 54 56 100
```

However, you can specify which columns you want to drop manually if you choose.

You'll usually have some justification for keeping or dropping some of them.

- Using a linear model and want a particular feature to serve as baseline
- Some dummies are too sparse

# findLinearCombos()

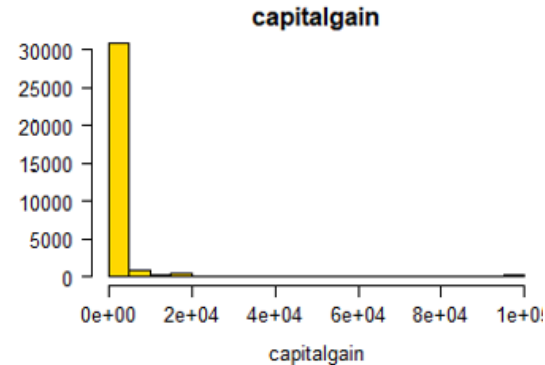

Lastly, we will go ahead and remove those columns indicated for removal. We also add back our target variable Y.

```
# remove columns identified that led to linear combos  
d <- d[, -comboInfo$remove]  
  
# remove the "ones" column in the first column  
d <- d[, c(2:ncol(d))]  
  
# Add the target variable back to our data.frame  
d <- cbind(y, d)  
  
rm(y, comboInfo) # clean up
```

# preProcess()

- 4) **preProcess()** function performs transformations (centering, scaling etc.) estimated from the training data and applied to any data set with the same variables

Variables tend to have ranges different from each other:



Some data mining algorithms are adversely affected by differences in variable ranges, where greater ranges tend to have larger influence on data model's results.

# preProcess()

There are various transformations available in preProcess. Some commonly used ones are:

**Goal:** Put your features on the same scale

- Z-score standardization
- Min-max normalization

```
method = c("center", "scale")  
method = c("range")
```

**Goal:** Make your features more bell-shaped

- Box-Cox
- Yeo-Johnson

```
method = c("BoxCox")  
method = c("YeoJohnson")
```

**Goal:** Do a combination of both

- Z-score & Yeo-Johnson
- Min-max & Yeo-Johnson

```
method = c("center", "scale", "YeoJohnson")  
method = c("range", "YeoJohnson")
```



# preProcess()

Here we transform our features using the min-max normalization (a.k.a. “range”)

```
#####  
# Standardize (and/ normalize) your input features.  
#####  
# Here we standardize the input features (Xs) using the preProcess() function  
# by performing a min-max normalization (aka "range" in caret).  
  
# Step 1) figures out the means, standard deviations, other parameters, etc. to  
# transform each variable  
preProcValues <- preProcess(d[,2:ncol(d)], method = c("range"))  
# Step 2) the predict() function actually does the transformation using the  
# parameters identified in the previous step. Weird that it uses predict() to do  
# this, but it does!  
d <- predict(preProcValues, d)
```

# preProcess()

Viewing the data post-transformation, you can see that the fnlwgt features is between 0 and 1.

	y	age	workclassGov	workclassPrivate	workclassSelf	fnlwgt	education10th	education11th	education12th	education1
1	<=50K	0.30136986	1	0	0	0.043337711	0	0	0	
2	<=50K	0.45205479	0	0	1	0.047277380	0	0	0	
3	<=50K	0.28767123	0	1	0	0.137243905	0	0	0	
4	<=50K	0.49315068	0	1	0	0.150211838	0	1	0	
5	<=50K	0.15068493	0	1	0	0.220703008	0	0	0	
6	<=50K	0.27397260	0	1	0	0.184109302	0	0	0	
7	<=50K	0.43835616	0	1	0	0.099540701	0	0	0	
8	>50K	0.47945205	0	0	1	0.133162150	0	0	0	

Previously, this feature was between 18 and 80+

# Couple things I do for classification probs...

Some algorithms wrapped in caret for classification require your target variable to be “named”. I always do this to avoid possible errors later in training.


```
#####  
# Get the target variable how we want it for modeling with caret  
#####  
# if greater than 50k make 1 other less than 50k make 0  
d$y <- as.factor(ifelse(d$y==" >50k",1,0))  
class(d$y)  
  
# make names for target if not already made  
levels(d$y) <- make.names(levels(factor(d$y)))  
levels(d$y)  
  
# levels of a factor are re-ordered so that the level specified is first and  
# "X1" is what we are predicting. The X before the 1 has nothing to do with the  
# X variables. It's just something weird with R. 'X1' is the same as 1 for the Y  
# variable and 'X0' is the same as 0 for the Y variable.  
d$y <- relevel(d$y,"X1")
```

# caret model design and training functions

- 5) **createDataPartition()** function allows one to easily partition their data into training and test sets that are distributed (or imbalanced) similar to one another.
- 6) **upSample()**, **downSample()** allows you to up or down sample your training data if it is severely unbalanced
- 7) **trainControl()** is where you specify how you want to design your run
- 8) **train()** is the bread and butter of what makes the caret package so great. It's a wrapper for essentially every (not all) regression or classification modeling technique

# createDataPartition()

5) **createDataPartition()** function allows one to easily partition their data into training and test sets that are distributed (or imbalanced) similar to one another.



# createDataPartition()

5) **createDataPartition()** function allows one to easily partition their data into training and test sets that are distributed (or imbalanced) similar to one another.

```
#####  
# Data partitioning  
#####  
set.seed(1234) # set a seed so you can replicate your results  
library(caret)  
  
# identify records that will be used in the training set. Here we are doing a  
# 70/30 train-test split. You might modify this.  
inTrain <- createDataPartition(y = d$y,      # outcome variable  
                               p = .70,      # % of training data you want  
                               list = F)  
  
# create your partitions  
train <- d[inTrain,] # training data set  
test  <- d[-inTrain,] # test data set
```

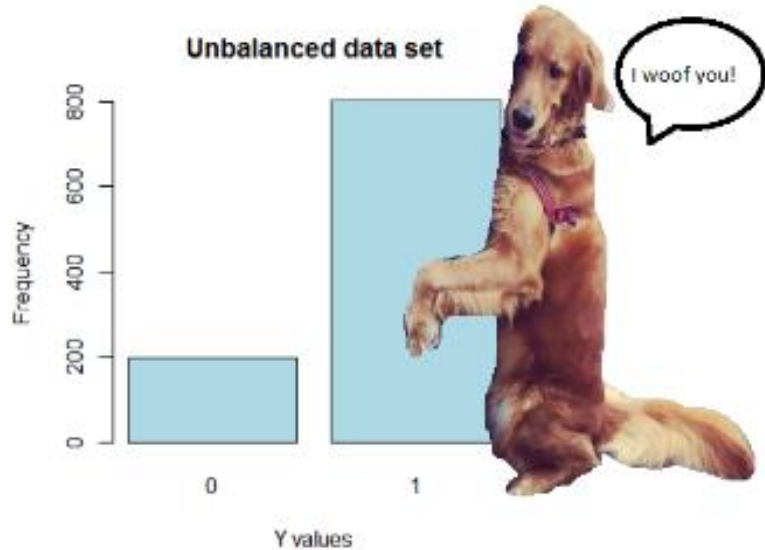
# createDataPartition()

```
#####  
# Data partitioning  
#####  
set.seed(1234) # set a seed so you can replicate your results  
library(caret)  
  
# identify records that will be used in the training set. Here we are doing a  
# 70/30 train-test split. You might modify this.  
inTrain <- createDataPartition(y = d$y,      # outcome variable  
                               p = .70,      # % of training data you want  
                               list = F)  
  
# create your partitions  
train <- d[inTrain,] # training data set  
test <- d[-inTrain,] # test data set
```

▶ test	9048 obs. of 91 variables
▶ train	21114 obs. of 91 variables

# upSample(), downSample()

6) upSample(), downSample() allows you to up or down sample your training data if it is severely unbalanced



Algorithms love the majority class like Golden Retrievers love tennis balls



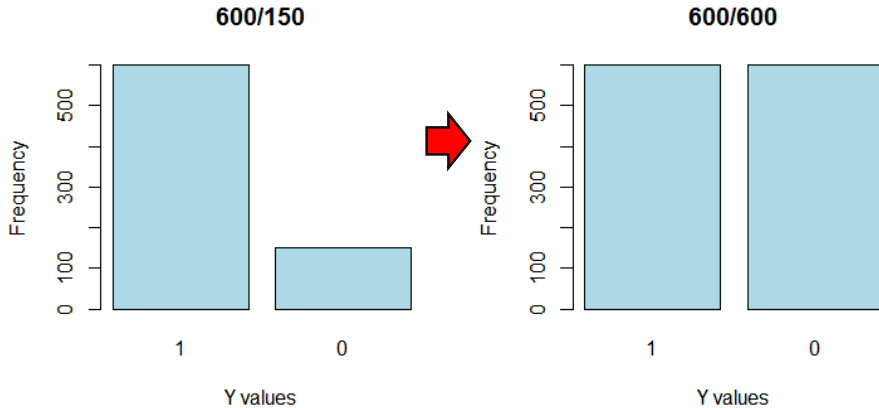


# upSample(), downSample()


Up-sampling resamples the minority class with replacement until the same number of records exist for the minority class as the majority class

Down-sampling samples the majority class until the same number of records exist for the majority class as the minority class

## Up-sampling



## Down-sampling



# upSample(), downSample()

Here is our down-sampled training set

```
# down-sampled training set  
dnTrain <- downSample(x=train[,2:ncol(d)], y=train$y)
```

▶ train	21114 obs. of 91 variables
▶ dnTrain	10512 obs. of 91 variables

# trainControl()

7) **trainControl()** is where you specify how you want to design your run

Here you specify if your problem is (1) a regression or (2) classification

Also, how you want to train your model (e.g. validation set approach, k-Fold cv, LOOCV, etc.)

```
#####  
# Specify cross-validation design  
#####  
ctrl <- trainControl(method="cv",      # cross-validation set approach to use  
                     number=3,        # k number of times to do k-fold  
                     classProbs = T,  # if you want probabilities  
                     summaryFunction = twoClassSummary, # for classification  
                     #summaryFunction = defaultSummary, # for regression  
                     allowParallel=T)
```

# trainControl()

7) **trainControl()** is where you specify how you want to design your run

Here you specify if your problem is (1) a regression or (2) classification

Also, how you want to train your model (e.g. validation set approach, k-Fold cv, LOOCV, etc.)

```
#####  
# Specify cross-validation design  
#####  
ctrl <- trainControl(method="cv",      # cross-validation set approach to use  
                     number=3,        # k number of times to do k-fold  
                     classProbs = T,  # if you want probabilities  
                     summaryFunction = twoClassSummary, # for classification  
                     #summaryFunction = defaultSummary, # for regression  
                     allowParallel=T)
```

# train()

8) **train()** is the bread and butter of what makes the caret package so great. It's a wrapper for essentially every (not all) regression or classification modeling technique





“Wrapper”?



# train() idea...



# train() idea...





# train() idea...



```
# you load Rose's library and use her custom function
library(Rose)
rose <- roseTree(x=Xmatrix, y=Y)

# you load Alice's library and use her custom function
library(Alice)
alice <- logitAlice(Y~X, data=myData)

# after you train the models, you want to generate predictions
predict(rose, type="prob")
predict(alice, type="raw")
```



No pre-specified design pattern so there are inconsistencies in how one sets functional inputs and what is returned in the output



# train() idea...

```
train <- function() {  
  library(Rose)  
  library(Alice)  
  library(Hongxia)  
  library(Shenyang)  
  
  # build Rose's decision tree  
  model <- roseTree(x=Xmatrix, y=Y)  
  
  # build Alice's logistic regression  
  model <- logitAlice(Y~X, data=myData)  
  
  # build Hongxia's SVM  
  model <- svmHongxia(Y~X, data=myData)  
  
  # build Shenyang's neural net  
  model <- nnShenyang(x=Xmatrix, y=Y, nodes=5)  
}
```

**train()** essentially wraps the packages (or writes a function around them) so there is a consistent way to input the arguments

# train()

All the available wrapped models for train are here:  
<https://topepo.github.io/caret/available-models.html>

## 6 Available Models

The models below are available in `train`. The code behind these protocols can be obtained using the function `getModelInfo` or by going to the [github repository](#).

Show  entries

Search:

Model	<i>method</i> Value	Type	Libraries
AdaBoost Classification Trees	adaboost	Classification	fastAdaboost
AdaBoost.M1	AdaBoost.M1	Classification	adabag, plyr
Adaptive Mixture Discriminant Analysis	amdai	Classification	adaptDA
Adaptive-Network-Based Fuzzy Inference System	ANFIS	Regression	frbs

# train()

Our first model is a logistic regression training on the train set (not down-sampled dataset)

Because glm is a family of different models, you need to specify which one (e.g. binomial/logit link) as an additional argument.

The models below are available in `train`. The code behind these protocols can be obtained using the function `getModelInfo` or by going to the [github repository](#).

Show  entries

Search:

Model	method Value	Type	Libraries	Tuning Parameters
Bayesian Generalized Linear Model	bayesglm	Classification, Regression	arm	None
Boosted Generalized Linear Model	glmboost	Classification, Regression	plyr, mboost	mstop, prune
Ensembles of Generalized Linear Models	randomGLM	Classification, Regression	randomGLM	maxInteractionOrder
Generalized Linear Model	glm	Classification, Regression		None

```
# train a logistic regression on train set
myModel1 <- train(y ~ .,
  data = train,
  method = "glm",
  trControl = ctrl,
  family = "binomial",
  metric = "ROC"
)
```

*# model specification*  
*# train set used to build model*  
*# type of model you want to build*  
*# how you want to learn*  
*# specify the type of glm*  
*# performance measure*

# train()

## Specify method

```
#####  
# Train different models  
#####  
# train a logistic regression on train set  
myModel1 <- train(y ~ .,                # model specification  
                  data = train,          # train set used to build model  
                  method = "glm",       # type of model you want to build  
                  trControl = ctrl,      # how you want to learn  
                  family = "binomial",   # specify the type of glm  
                  metric = "ROC"        # performance measure  
)  
myModel1  
  
# train a logistic regression on down-sampled train set  
myModel2 <- train(y ~ .,                # model specification  
                  data = dnTrain,        # train set used to build model  
                  method = "glm",       # type of model you want to build  
                  trControl = ctrl,      # how you want to learn  
                  family = "binomial",   # specify the type of glm  
                  metric = "ROC"        # performance measure  
)  
myModel2
```

# train()

Maybe we'll try a feed-forward neural net.

The models below are available in `train`. The code behind these protocols can be obtained using the function `getModelInfo` or by going to the [github repository](#).

Show  entries

Search:

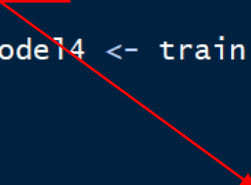
Model	<i>method</i> Value	Type	Libraries	Tuning Parameters
Model Averaged Neural Network	avNNet	Classification, Regression	nnet	size, decay, bag
Neural Network	nnet	Classification, Regression	nnet	size, decay

```
# train a feed-forward neural net on train set
myModel13 <- train(y ~ .,
  data = train,
  method = "nnet",
  trControl = ctrl,
  tuneLength = 1:5,
  maxit = 100,
  metric = "ROC"
)
myModel13
```

*# model specification*  
*# train set used to build model*  
*# type of model you want to build*  
*# how you want to learn*  
*# how many tuning parameter combos to try*  
*# max # of iterations*  
*# performance measure*

# train()

```
# train a feed-forward neural net on train set
myModel3 <- train(y ~ .,                                # model specification
                  data = train,                        # train set used to build model
                  method = "nnet",                    # type of model you want to build
                  trControl = ctrl,                    # how you want to learn
                  tuneLength = 1:5,                   # how many tuning parameter combos to try
                  maxit = 100,                        # max # of iterations
                  metric = "ROC"                      # performance measure
)
myModel3
|
# train a feed-forward neural net on the down-sampled train set using a customer
# tuning parameter grid
myGrid <- expand.grid(size = c(10,15,20)               # number of units in the hidden layer
                     , decay = c(.09,0.12))          # parameter for weight decay. Default
myModel4 <- train(y ~ .,                               # model specification
                  data = dnTrain,                     # train set used to build model
                  method = "nnet",                    # type of model you want to build
                  trControl = ctrl,                    # how you want to learn
                  tuneGrid = myGrid,                   # tuning parameter combos to try
                  maxit = 100,                        # max # of iterations
                  metric = "ROC"                      # performance measure
)
myModel4
```



# train()

```
> myModel4
Neural Network

10512 samples
  90 predictor
  2 classes: 'x1', 'x0'

No pre-processing
Resampling: Cross-Validated (3 fold)
Summary of sample sizes: 7008, 7008, 7008
Resampling results across tuning parameters:
```

size	decay	ROC	Sens	Spec
10	0.09	0.8925994	0.8424658	0.7886225
10	0.12	0.8949109	0.8390411	0.7792998
15	0.09	0.0000000	NaN	NaN
15	0.12	0.0000000	NaN	NaN
20	0.09	0.0000000	NaN	NaN
20	0.12	0.0000000	NaN	NaN

ROC was used to select the optimal model using the largest value.  
The final values used for the model were size = 10 and decay = 0.12.

Shows the best set of tuning parameters at the bottom.

# caret evaluation functions

- 9) **predict()** allows you to generate your predictions from a training set or testing set
- 10) **confusionMatrix()** provides a convenient function to assess classification model performance



# predict()

9) `predict()` allows you to generate your predictions from a training set or testing set

```
# Capture the train and test estimated probabilities and predicted classes  
# model 1  
logit1_trp <- predict(myModel1, newdata=train, type='prob')[,1]  
logit1_trc <- predict(myModel1, newdata=train)  
logit1_tep <- predict(myModel1, newdata=test, type='prob')[,1]  
logit1_tec <- predict(myModel1, newdata=test)  
# model 2  
logit2_trp <- predict(myModel2, newdata=dnTrain, type='prob')[,1]  
logit2_trc <- predict(myModel2, newdata=dnTrain)  
logit2_tep <- predict(myModel2, newdata=test, type='prob')[,1]  
logit2_tec <- predict(myModel2, newdata=test)  
# model 3  
nn1_trp <- predict(myModel3, newdata=train, type='prob')[,1]  
nn1_trc <- predict(myModel3, newdata=train)  
nn1_tep <- predict(myModel3, newdata=test, type='prob')[,1]  
nn1_tec <- predict(myModel3, newdata=test)  
# model 4  
nn2_trp <- predict(myModel4, newdata=dnTrain, type='prob')[,1]  
nn2_trc <- predict(myModel4, newdata=dnTrain)  
nn2_tep <- predict(myModel4, newdata=test, type='prob')[,1]  
nn2_tec <- predict(myModel4, newdata=test)
```

# confusionMatrix()

**10) confusionMatrix()** provides a convenient function to assess classification model performance

```
# Now use those predictions to assess performance on the train set and testing  
# set. Be on the lookout for overfitting  
# model 1 - logit  
(cm <- confusionMatrix(data=logit1_trc, train$y))  
(testCM <- confusionMatrix(data=logit1_tec, test$y))  
# model 2 - logit with down-sampled data  
(cm2 <- confusionMatrix(data=logit2_trc, dnTrain$y))  
(testCM2 <- confusionMatrix(data=logit2_tec, test$y))  
# model 3 - nnet  
(cm3 <- confusionMatrix(data=nn1_trc, train$y))  
(testCM3 <- confusionMatrix(data=nn1_tec, test$y))  
# model 4 - nnet with down-sampled data  
(cm4 <- confusionMatrix(data=nn2_trc, dnTrain$y))  
(testCM4 <- confusionMatrix(data=nn2_tec, test$y))
```

# confusionMatrix()

Neural net model without rebalancing

training

```
> (cm3 <- confusionMatrix(data=nn1_trc, train$y))  
Confusion Matrix and Statistics
```

	Reference	
Prediction	X1	X0
X1	3361	1326
X0	1895	14532

Accuracy : 0.8474  
95% CI : (0.8425, 0.8523)  
No Information Rate : 0.7511  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.5767  
McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.6395  
Specificity : 0.9164  
Pos Pred Value : 0.7171  
Neg Pred Value : 0.8846  
Prevalence : 0.2489  
Detection Rate : 0.1592  
Detection Prevalence : 0.2220  
Balanced Accuracy : 0.7779

'Positive' Class : X1

test

```
> (testCM3 <- confusionMatrix(data=nn1_tec, test$y))  
Confusion Matrix and Statistics
```

	Reference	
Prediction	X1	X0
X1	1438	594
X0	814	6202

Accuracy : 0.8444  
95% CI : (0.8368, 0.8518)  
No Information Rate : 0.7511  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.5697  
McNemar's Test P-Value : 5.335e-09

Sensitivity : 0.6385  
Specificity : 0.9126  
Pos Pred Value : 0.7077  
Neg Pred Value : 0.8840  
Prevalence : 0.2489  
Detection Rate : 0.1589  
Detection Prevalence : 0.2246  
Balanced Accuracy : 0.7756

'Positive' Class : X1

# Take Away

- **caret provides**
  - A nice wrapper to many (not all) R packages for predictive modeling and machine learning
  - Functions to easily pre-process your data before modeling
  - Ability to tune hyperparameters with ease by specifying **tuneLength=** or **tuneGrid=**
  - Can incorporate other specific package arguments, but they might not be part of tuning parameters list
  - Easily evaluate regression and classification-type problems